

# Apache Metron

## A Case Study of a Modern Streaming Architecture on Hadoop

**Casey Stella**  
@casey\_stella



2017

## Introduction

---

Hi, I'm Casey Stella!

# Apache Metron: A Cybersecurity Analytics Platform

---

- Metron provides a scalable, advanced security analytics framework to offer a centralized tool for security monitoring and analysis.
- Metron was initiated at Cisco in 2014 as OpenSOC.
- Metron was submitted to the Apache Incubator in December 2015
- Metron graduated to a top level project in April 2017

# Characteristics of Metron

---

- Metron is built atop the Apache Hadoop ecosystem handle capturing, ingesting, enriching and storing streaming data at scale
  - Kafka provides a unified data bus
  - Storm providing a distributed streaming framework
  - HBase provides a low latency key/value lookup store for enrichments and profiles
  - Zookeeper provides a distributed configuration store
- Ingested network telemetry can be enriched pluggably
  - New enrichments can be done live on running topologies without restart
  - New enrichment capabilities can be added via user defined functions
  - Enrichments can be composed through a domain specific language called **Stellar**
- Data stored in HBase can be the source of enrichments

# Characteristics of Metron

---

- Enriched network telemetry can be indexed into a Security data lake
  - Indexes supported are pluggable and include HDFS, Solr and Elasticsearch
- Advanced analytics can be done on streaming data
  - Probabalistic data structures (e.g. sketches) can sketch streaming data across time and enable approximate distribution, set existence and distinct count queries
  - Models can be deployed using Yarn, autodiscovered via Zookeeper and interrogated via Stellar functions

# Stellar

---

Metron needed the ability to allow users to pluggably and consistently enrich and transform streaming data. Out of this need, we created **Stellar**:

- Interact with the various enabling Hadoop components in a unified manner
- Compose a rich set of built-in functions with user defined functions
- Provide simple primitives around the functions: boolean operations, conditionals, numerical computation.


Think of Stellar as similar to Excel functions that we can run on streaming data.

## Data Ingest: Parsers

---

- Telemetry data comes in a variety of formats and velocities.
- Each telemetry source is ingested into kafka
- A Storm **parser** topology is used to convert the raw telemetry format to a normalized JSON Map
  - Common network-related raw telemetry formats
  - Specifying the parser via Grok
  - Generic formats such as CSV and JSON
  - Creating your own parser in a JVM-based language
- Simple transformations and normalizations can be done post-parse via Stellar statements
- The normalized data across all telemetries is written to an **enrichment** kafka topic

# Data Ingest: Parsers





# Enrichment

---

- The enrichment topology takes the various normalized telemetry sources and allows users to enrich the messages with broader context
  - Enriching with reference data ingested into HBase
  - Enriching via arbitrary Stellar expressions
  - Enriching with Geolocation data
- Enrichment is split into two phases
  - Preparatory Enrichment
  - Threat Intelligence Enrichment
- If messages are marked with an **is\_alert** field, they can have a triage score computed via Stellar which defines their priority as threats

## Enrichment: HBase

---

A core enrichment source is enriching messages with reference data stored in HBase. Metron provides a loading framework which takes data and

- Defines a logical key for the HBase enrichment via Stellar
- Loads the data into HBase via MapReduce or directly depending on data size
- For reference data which is streamed in, you can write to HBase from a parser and ingest the streaming reference data as any other telemetry

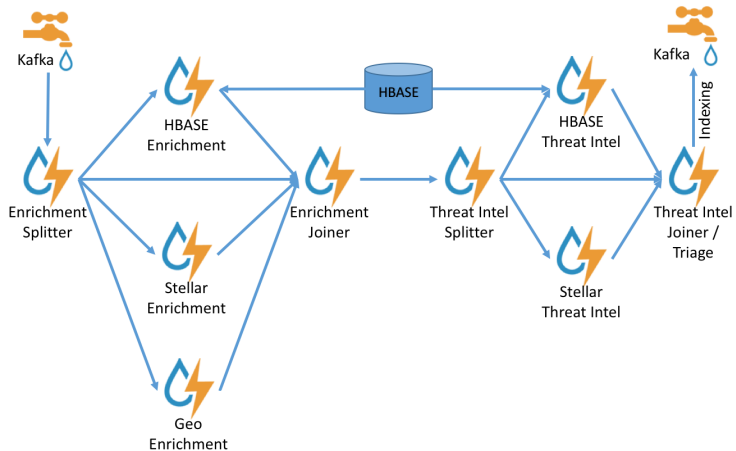
## Enrichment: Stellar

---


Stellar is the primary method for enrichment in Metron.

- User defined enrichment functions can be enabled through adding a jar implementing the function to HDFS.
- Stellar enrichments can be executed asynchronously across storm workers and their results joined together

# Enrichment



# Indexing



## Profiler: Motivation

---

- Enrichments and parsers operate within the context of a single message.
- This is insufficient for a number of scenarios
  - Correlating between different sources
  - Making judgments based on past events
  - Both at the same time
- Operating across multiple sources has scalability implications
- Waiting on the data you want from the other stream isn't plausible

## Profiler: Solution

---

- Compromise: Operate on windows in time rather than individual records
- Windows should be able to be specified very flexibly to avoid seasonal aberrations.
- The Profiler is a storm topology that takes the enriched data
  - Capture aggregations of data in a window to HBase
  - Uses Stellar to define how to aggregate data
  - Uses Stellar to define a filter on which messages to consider
- These aggregations can be read anywhere Stellar is used
- This enables things like
  - Temporal outlier detection
  - Limited context from other sources when building threat triage rules

# Profiler: Aggregations

---

- Example aggregations:
  - Distributional statistics: median, mean, percentile
  - Set operations: Contains, Cardinality
  - Simple counts
- Aggregations challenges
  - May be big objects if naively done (e.g. set operations)
  - May not be able to be merged across time (e.g. distributions)
  - Profile reading should be decoupled from writing
- Approach: Use approximate data structures
  - Set operations: Bloom Filters, HyperLogLog approximations
  - Distributional Statistics: t-digests



# Demo

---

- Los Alamos National Lab released an open data set representing 58 consecutive days of de-identified event data collected from five sources within Los Alamos National Laboratory's corporate, internal computer network.
- Among other telemetry sources, authentication logs and a set of well-defined red teaming events that present bad behavior within the 58 days are provided.
- We will look at the authentication logs around a breach and show how we can use Metron to pick out offending user's activity leading up to the event
  - Look for users who are attempting to authenticate to many distinct hosts more than 5 standard deviations from the median across all users.

```
"profile": "distinct_auth_attempts_by_user",
"foreach": "user",
"onlyif": "source.type == 'auth'",
"init" : {
    "total" : "HLLP_INIT(5,6)"
},
"update": {
    "total" : "HLLP_ADD(total, ip_dst_addr)"
},
"result" : {
    "profile" : "total",
    "trriage" : {
        "total_count" : "HLLP_CARDINALITY(total)"
    }
}
```

```
"profile": "auth_distribution",
"foreach": "'global'",
"onlyif": "source.type == 'profiler' && profile == 'distinct_auth_attempts_by_us",
"init" : {
    "s" : "STATS_INIT()"
    },
"update": {
    "s" : "STATS_ADD(s, total_count)"
    },
"result": "s"
```

```
window := PROFILE_WINDOW('...')
profile := PROFILE_GET('distinct_auth_attempts_by_user', user, window)
distinct_auth_attempts := HLLP_CARDINALITY(GET_LAST(profile))
distribution_profile := PROFILE_GET('auth_distribution', 'global', window)
stats := STATS_MERGE(distribution_profile)
distinct_auth_attempts_median := STATS_PERCENTILE(stats, 0.5)
distinct_auth_attempts_stddev := STATS_SD(stats)
```

# Questions

---

Thanks for your attention! Questions?

- Find me at <http://caseystella.com>
- Twitter handle: @casey\_stella
- Email address: [cstella@hortonworks.com](mailto:cstella@hortonworks.com)